

Boot-time

```
main.c

// Bootstrap processor starts running C code here.

int main(void)
{
    // ...
    tvinit();           // trap vectors
}
```

```
traps.h

Interrupt name->number mappings. Name is used
elsewhere in code, rather than number.

// Defined by Intel, table 6-1, page 152
T_DIVIDE    0 // divide error
T_DEBUG     1 // debug exception
T_NMI       2 // non-maskable interrupt
T_BRKPT     3 // breakpoint
T_OFLOW     4 // overflow
T_BOUND     5 // bounds check
T_ILLOP     6 // illegal opcode
T_DEVICE    7 // device not available
T_DBLFLT    8 // double fault
T_TSS       10 // invalid task switch segment
T_SEGNP     11 // segment not present
T_STACK     12 // stack exception
T_GPFLT     13 // general protection fault
T_PGFLT     14 // page fault
T_FPERR     16 // floating point error
T_ALIGN     17 // alignment check
T_MCHK      18 // machine check
T_SIMDERR   19 // SIMD floating point error

// xv6-specific interrupt mappings
T_SYSCALL  64 // system call
T_DEFAULT   500 // catchall

T_IRQ0     32 // IRQ 0
IRQ_TIMER  0 // timer interrupt
```

```
syscall.h

// System call numbers
#define SYS_fork 1 // cf. 'num' in syscall.c
#define SYS_exit 2
// ...
```

```
trap.c

static struct gatedesc idt[256];

void tvinit()
{
    for(i = 0; i < 256; i++)
        SETGATE(idt[i], 0, SEG_KCODE<<3, vectors[i], 0);
    SETGATE(idt[T_SYSCALL], 1, SEG_KCODE<<3, vectors[T_SYSCALL], DPL_USER);
}

void idtinit(void)
{
    lidt(idt, sizeof(idt)); // LIDT ASM instruction loads IDT table
}

mmu.h

// Set up a normal interrupt/trap gate descriptor.
// - istrap: 1 for a trap (= exception) gate, 0 for an interrupt gate.
// - interrupt gate clears FL_IF, trap gate leaves FL_IF alone
// - sel: Code segment selector for interrupt/trap handler
// - off: Offset in code segment for interrupt/trap handler
// - dpl: Descriptor Privilege Level -
//       the privilege level required for software to invoke
//       this interrupt/trap gate explicitly using an int instruction.
#define SETGATE(gate, istrap, sel, off, d)
{
    (gate).off_15_0 = (uint)(off) & 0xffff;
    (gate).cs = (sel);
    (gate).args = 0;
    (gate).rsv1 = 0;
    (gate).type = (istrap) ? STS_TG32 : STS_IG32;
    (gate).s = 0;
    (gate).dpl = (d);
    (gate).p = 1;
    (gate).off_31_16 = (uint)(off) >> 16;
}

void trap(struct trapframe *tf)
{
    if(tf->trapno == T_SYSCALL){
        if(proc->killed)
            exit();
        proc->tf = tf;
        syscall();
        if(proc->killed)
            exit();
        return;
    }
    // ...
    if(proc &&
        proc->state == RUNNING &&
        tf->trapno == T_IRQ0+IRQ_TIMER)
        yield();
}
```

Run-time

```
User program (ASM code)

// ... push relevant arguments on stack
call fork
```

```
usys.S

#define SYSCALL(name) \ // name is, e.g., 'fork'
    .globl name; \
    name: \
        // turns into, e.g., 'movl $SYS_fork, %eax;'
        movl $SYS_ ## name, %eax; \
        int T_SYSCALL; \ // cause specific interrupt
        ret

SYSCALL(fork)
SYSCALL(exit)
// ...
```

```
trapasm.S

.globl alltraps
alltraps:
    # Build trap frame.
    pushl %ds
    pushl %es
    pushl %fs
    pushl %gs
    pushal
    //...
    # Call trap(tf), where tf=%esp
    pushl %esp
    call trap
```

```
vectors.S

# generated by vectors.pl
# do not edit
# handlers
.globl alltraps
.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps
.globl vector1
vector1:
    pushl $0
    pushl $1
    jmp alltraps
// ...
# vector table
.data
.globl vectors
vectors:
    .long vector0
    .long vector1
// ...
```

```
syscall.c

extern int sys_fork(void);
extern int sys_exit(void);

static int (*syscalls[])(void) = {
    [SYS_fork] = sys_fork,
    [SYS_exit] = sys_exit,
    // ...
};

void syscall(void)
{
    int num;

    num = proc->tf->eax;
    if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
        proc->tf->eax = syscalls[num](); // set eax to return value
    } else {
        cprintf("%d %s: unknown sys call %d\n",
            proc->pid, proc->name, num);
        proc->tf->eax = -1;
    }
}
```

```
proc.c

int sys_fork(void)
{
    return fork();
}

int fork(void)
{
    // ... actual fork code
    return pid;
}
```